

Introduction to Vulkan


A High Level Walk Through Modern Graphics
API Concepts

Dan Ginsburg – Upsample Software, LLC

Overview

- ▶ About Me
- ▶ A Bit of Graphics APIs History
- ▶ Why Vulkan?
- ▶ Key Modern Graphics API Concepts
- ▶ Vulkan Resources

Stuff I've worked on...



Many possible career paths for graphics programmers...

- ▶ Drivers
 - ▶ **NVIDIA, AMD, Qualcomm, Intel, ARM, Apple, ImgTec, etc.**
- ▶ GPU Tools, GPU Research
 - ▶ **Same companies as above, + game engines/devs**
- ▶ Rendering
 - ▶ **Game developers/engines**
- ▶ Many other fields that use graphics/GPUs
 - ▶ **Medical Devices**
 - ▶ **Avionics/Automotive**
 - ▶ **User Interfaces**
 - ▶ **CAD**
 - ▶ ...

A Bit of Graphics API History

How did WebGL come to be?

Let's go back 15 years...2006



What is OpenGL ES?

- OpenGL-based API for embedded systems
- Removed redundant and expensive functionality
- Kept as compatible with OpenGL as possible
- Added features needed for embedded systems

What is OpenGL ES?

- OpenGL-based API for embedded systems
- Removed redundant and expensive functionality
- Kept as compatible with OpenGL as possible
- Added features needed for embedded systems

OpenGL ES History

- OpenGL ES 1.0
 - Basic 3D functionality
- OpenGL ES 1.1+
 - Comprehensive set of fixed-function 3D
 - Backwards compatible
- OpenGL ES 2.0
 - Shader only
 - Not backwards compatible

iPhone 3G S – June 2008



OpenGL ES and WebGL Evolution

Pervasive OpenGL ES 2.0

OpenGL and OpenGL ES ships on every desktop and mobile OS
3D on the Web is enabled!

Mobile Graphics

Programmable Vertex and Fragment shaders



2007

OpenGL ES 2.0

Desktop Graphics

Textures: NPOT, 3D, Depth, Arrays, Int/float
Objects: Query, Sync, Samplers
Seamless Cubemaps, Integer vertex attributes
Multiple Render Targets, Instanced rendering
Transform feedback, Uniform blocks
Vertex array objects, GLSL ES 3.0 shaders



2012

OpenGL ES 3.0



Compute Shaders



2014

OpenGL ES 3.1
Compute Shaders

Advanced Graphics

Tessellation and geometry shaders
ASTC Texture Compression
Floating point render targets
Debug and robustness for security



2015

OpenGL ES 3.2



4 years

2011
WebGL 1.0

WebGL 2.0 Compute Context
Multiview extension

Work in Progress

5 years

March 2017
WebGL 2.0


Conformance Testing is vital for Cross-Platform Reliability

WebGL 2.0 conformance tests are very thorough 10x more tests than WebGL 1.0 tests

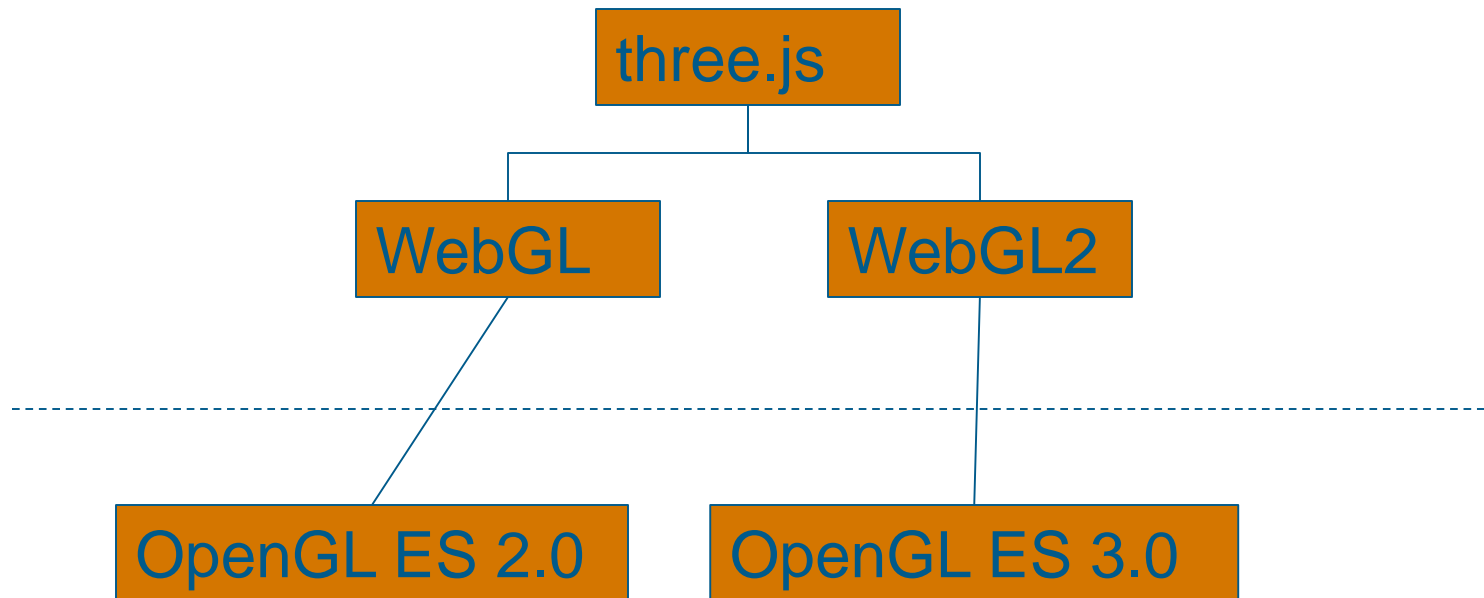
Source:

https://www.web3d.org/sites/default/files/attachment/node/2326/edit/49_HwanyongLee_KhronosLiaisonReport_Seoul2019January24.pdf

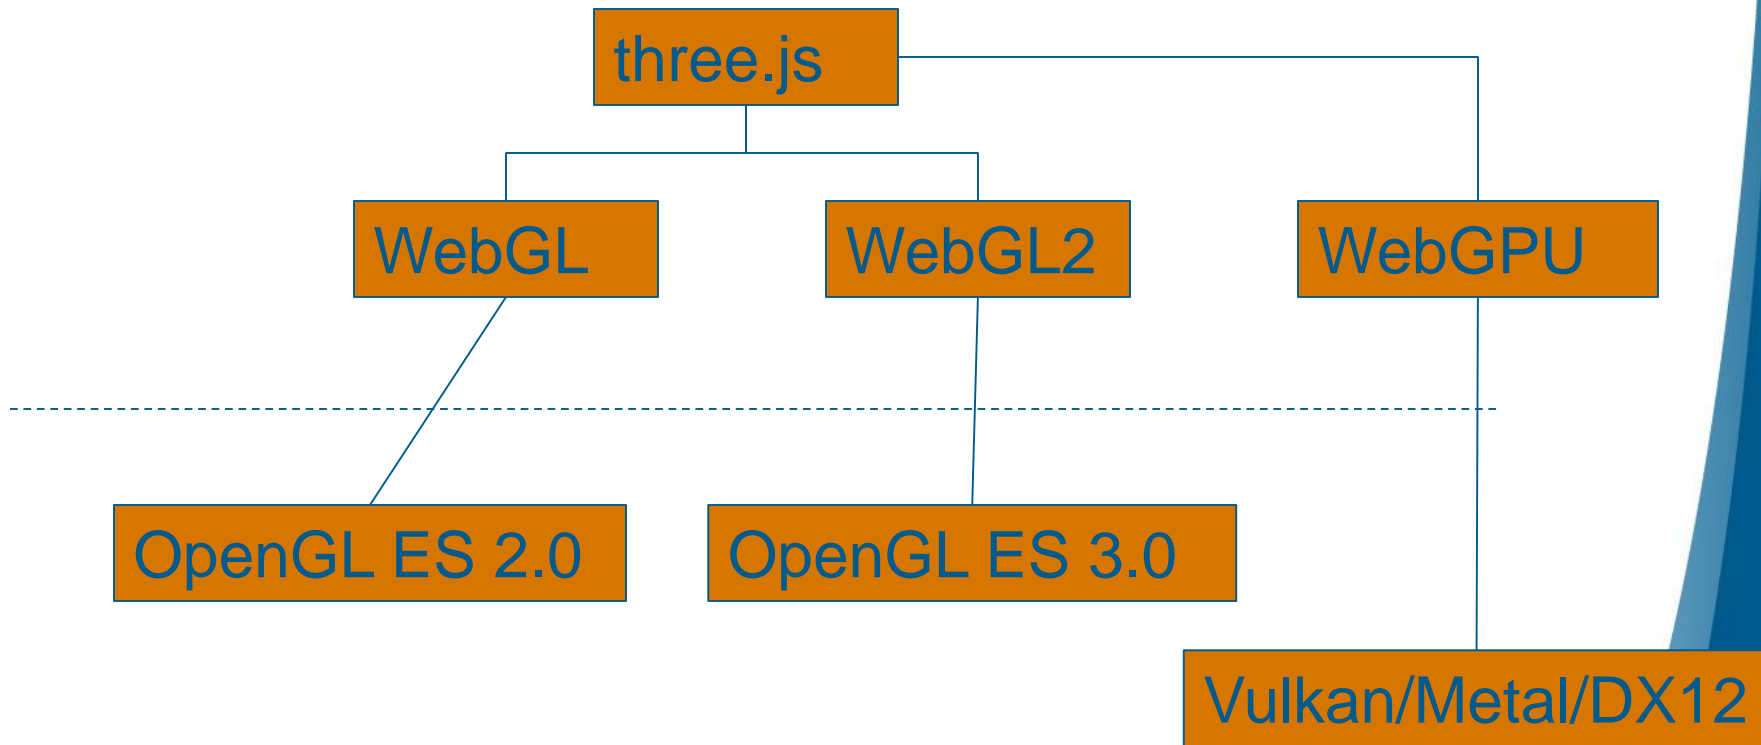
three.js Web APIs



three.js Web APIs




three.js Web APIs Future

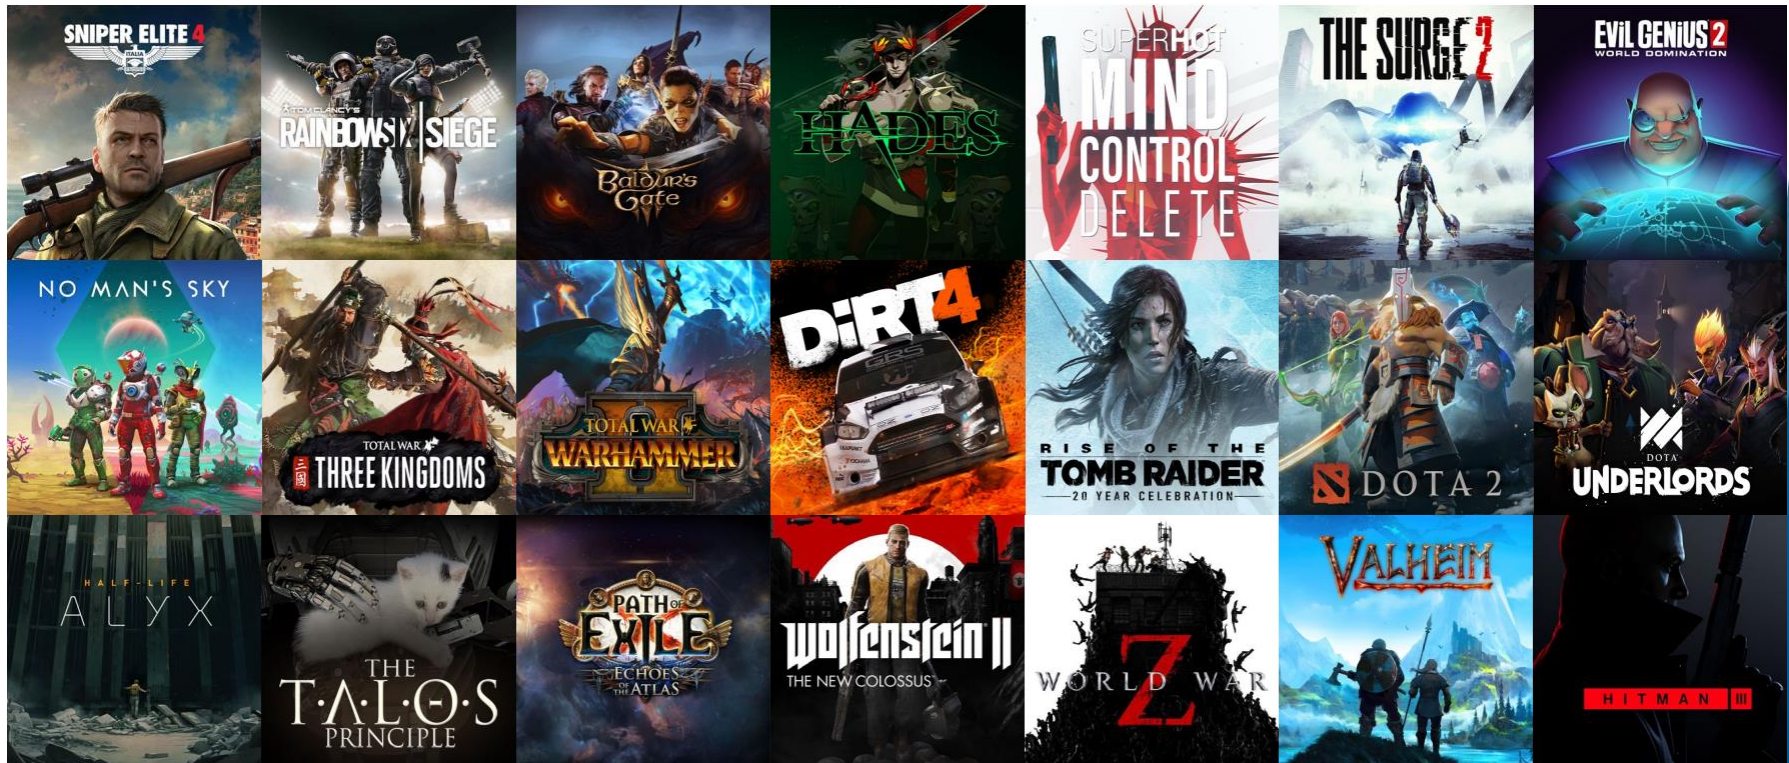


Chrome WebGL Implementation

- ▶ ANGLE: <https://github.com/google/angle>



Why Vulkan?



OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API




OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



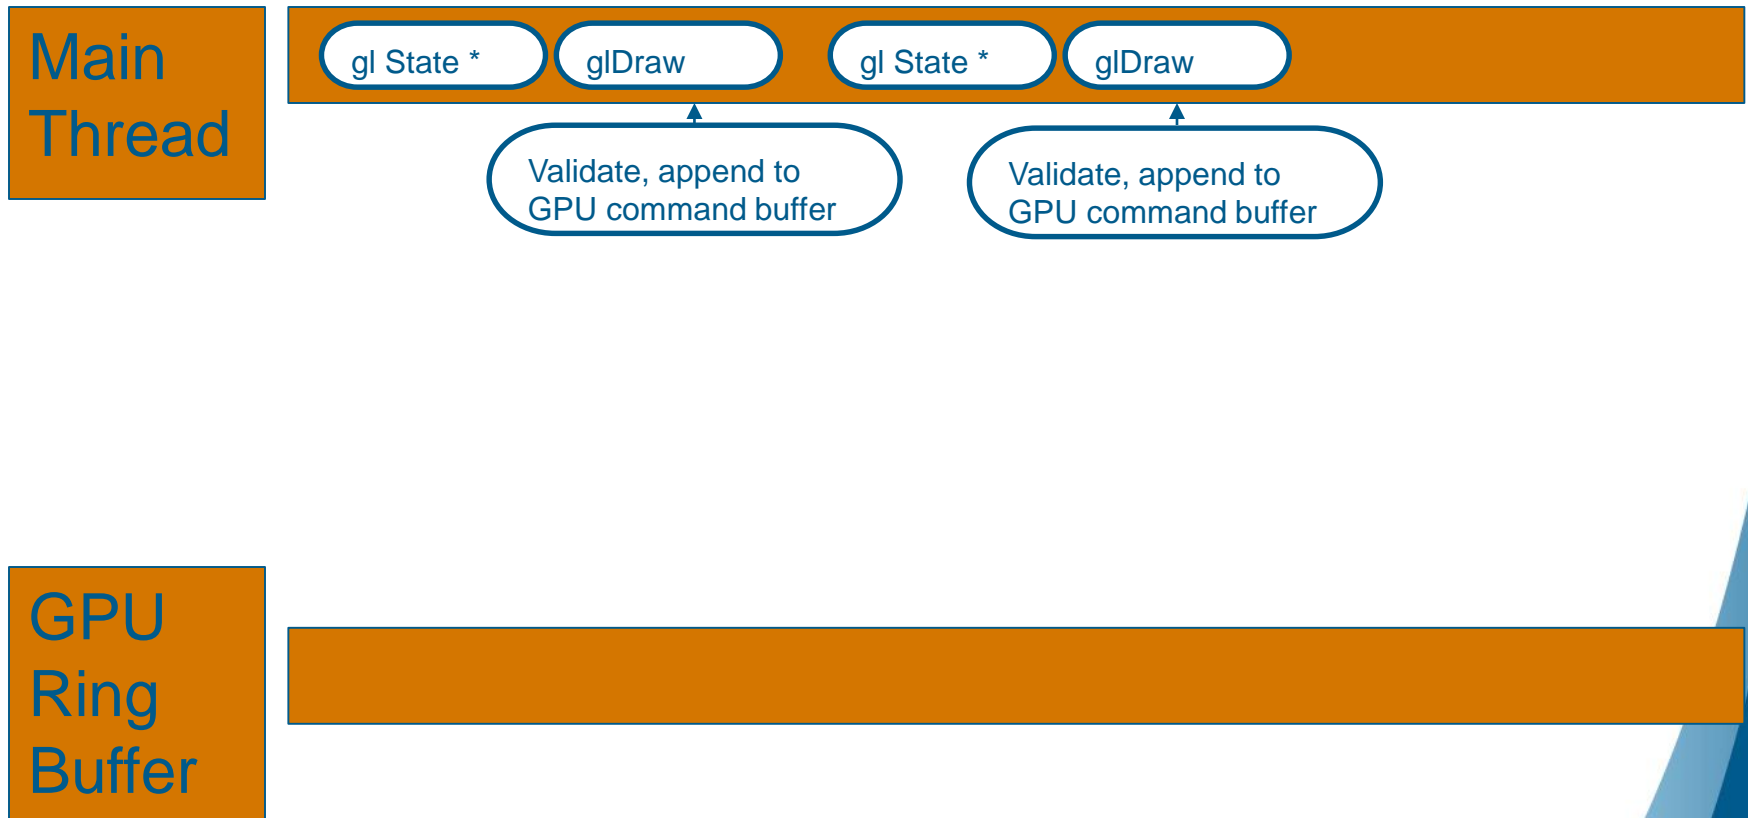
OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



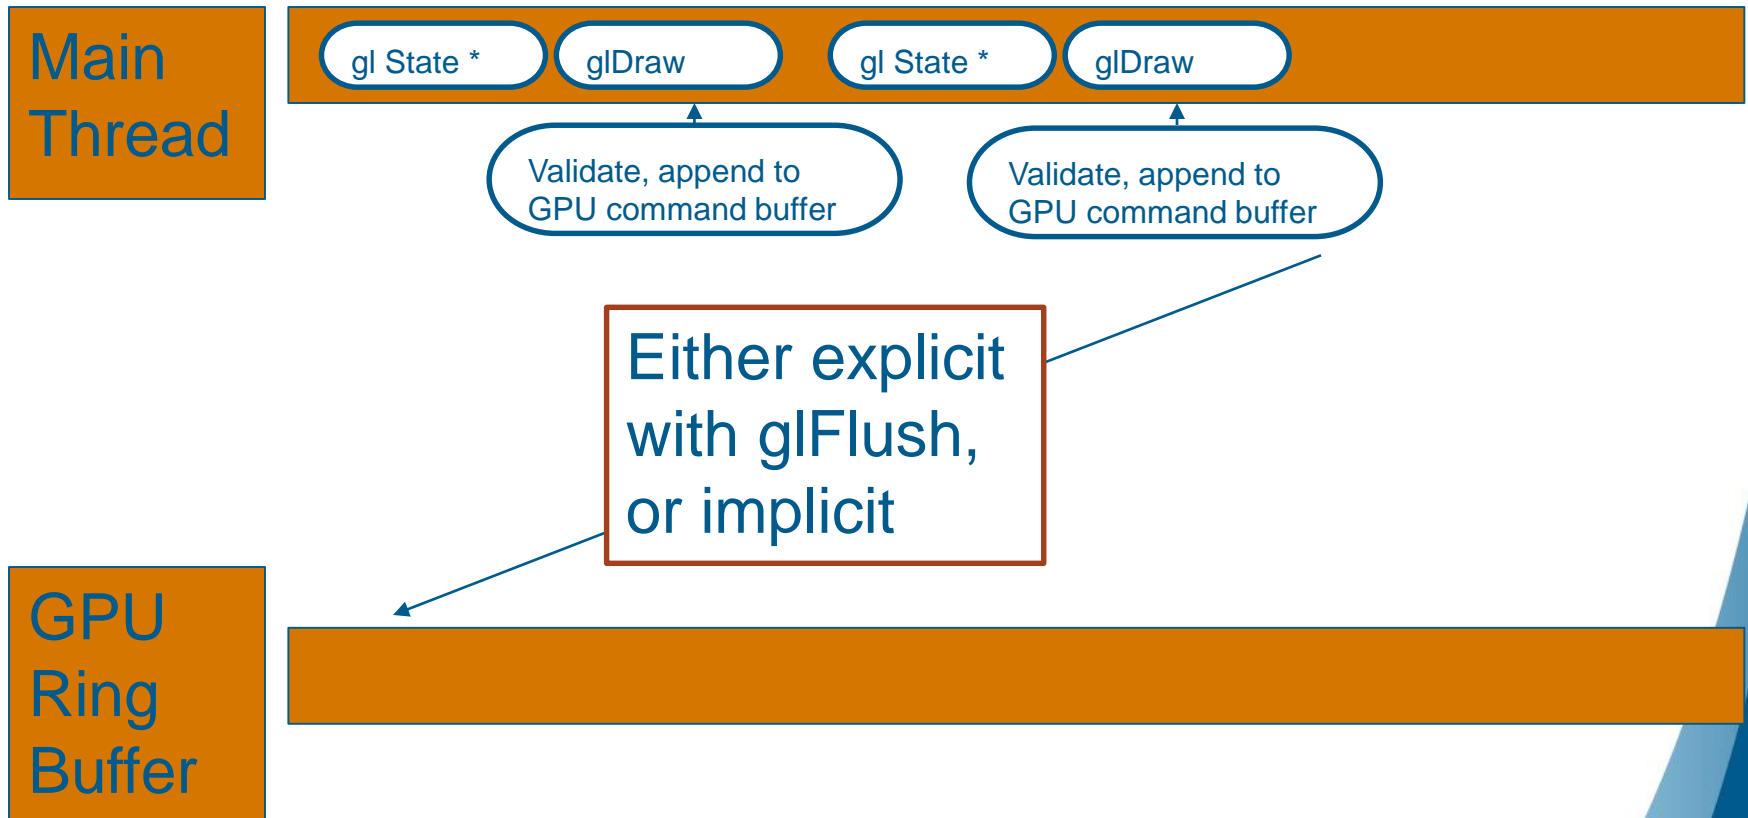
OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



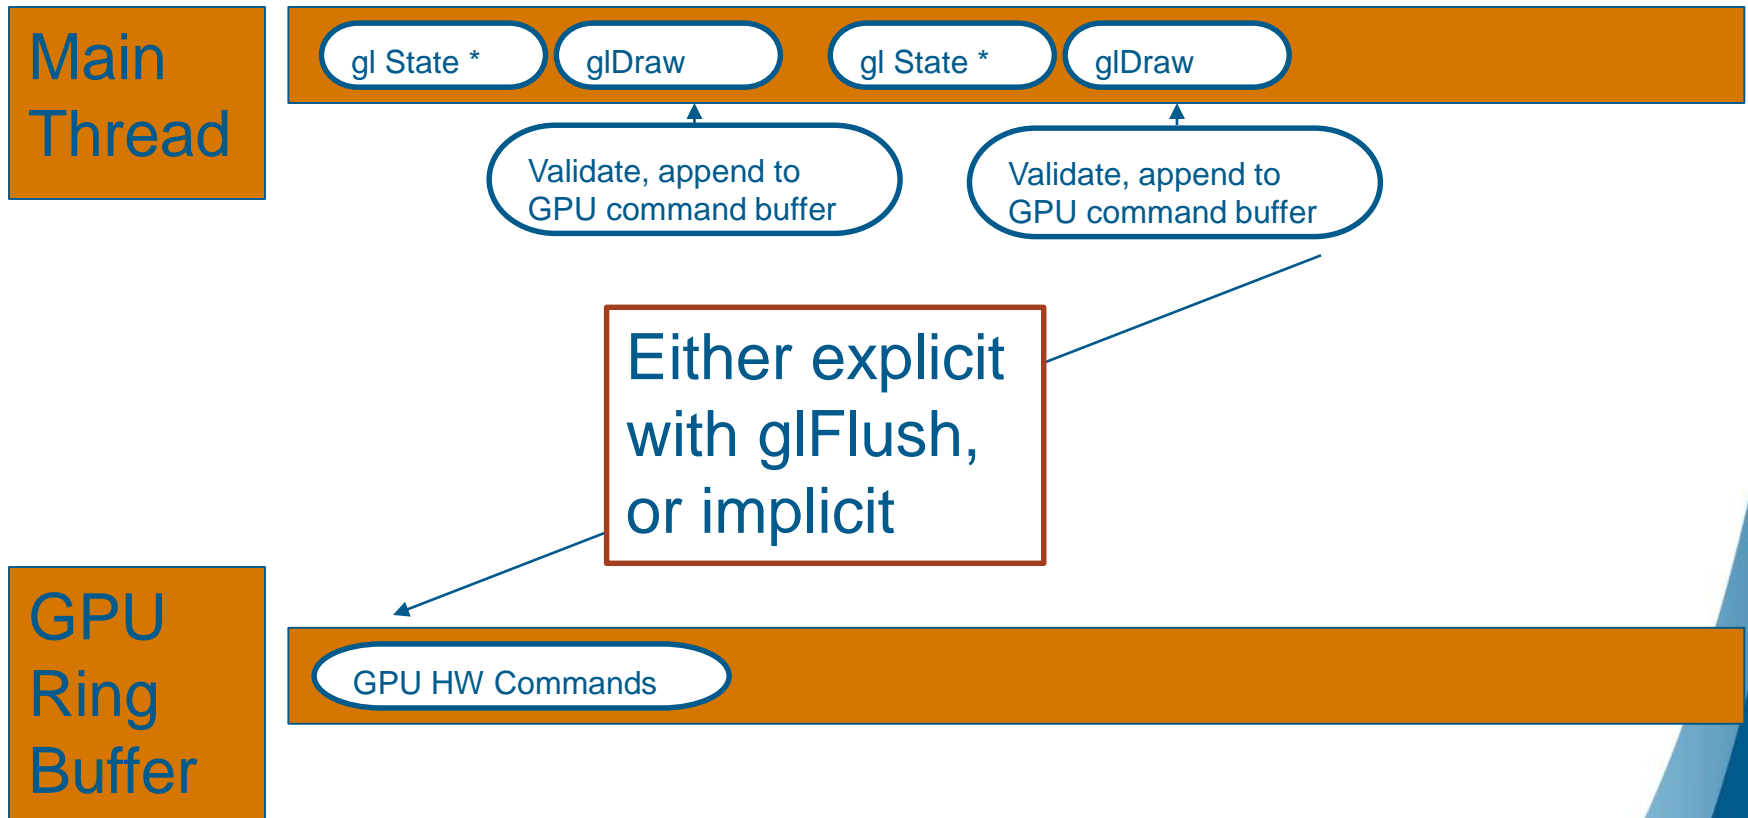
OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



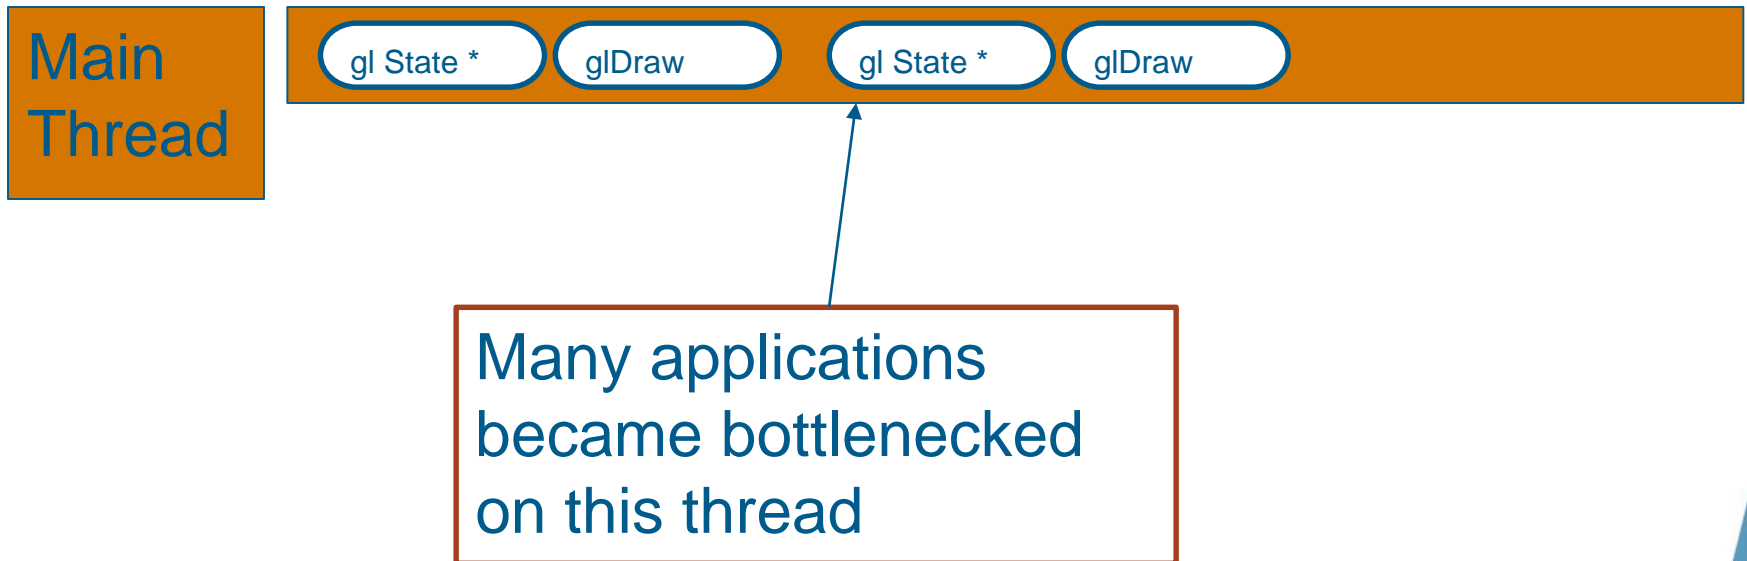
OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



OpenGL/OpenGL ES Submission Model

- ▶ One Thread Submits to the API



Multithreading OpenGL/OpenGLES

Thread
A

SW Command Buffer

Thread
B

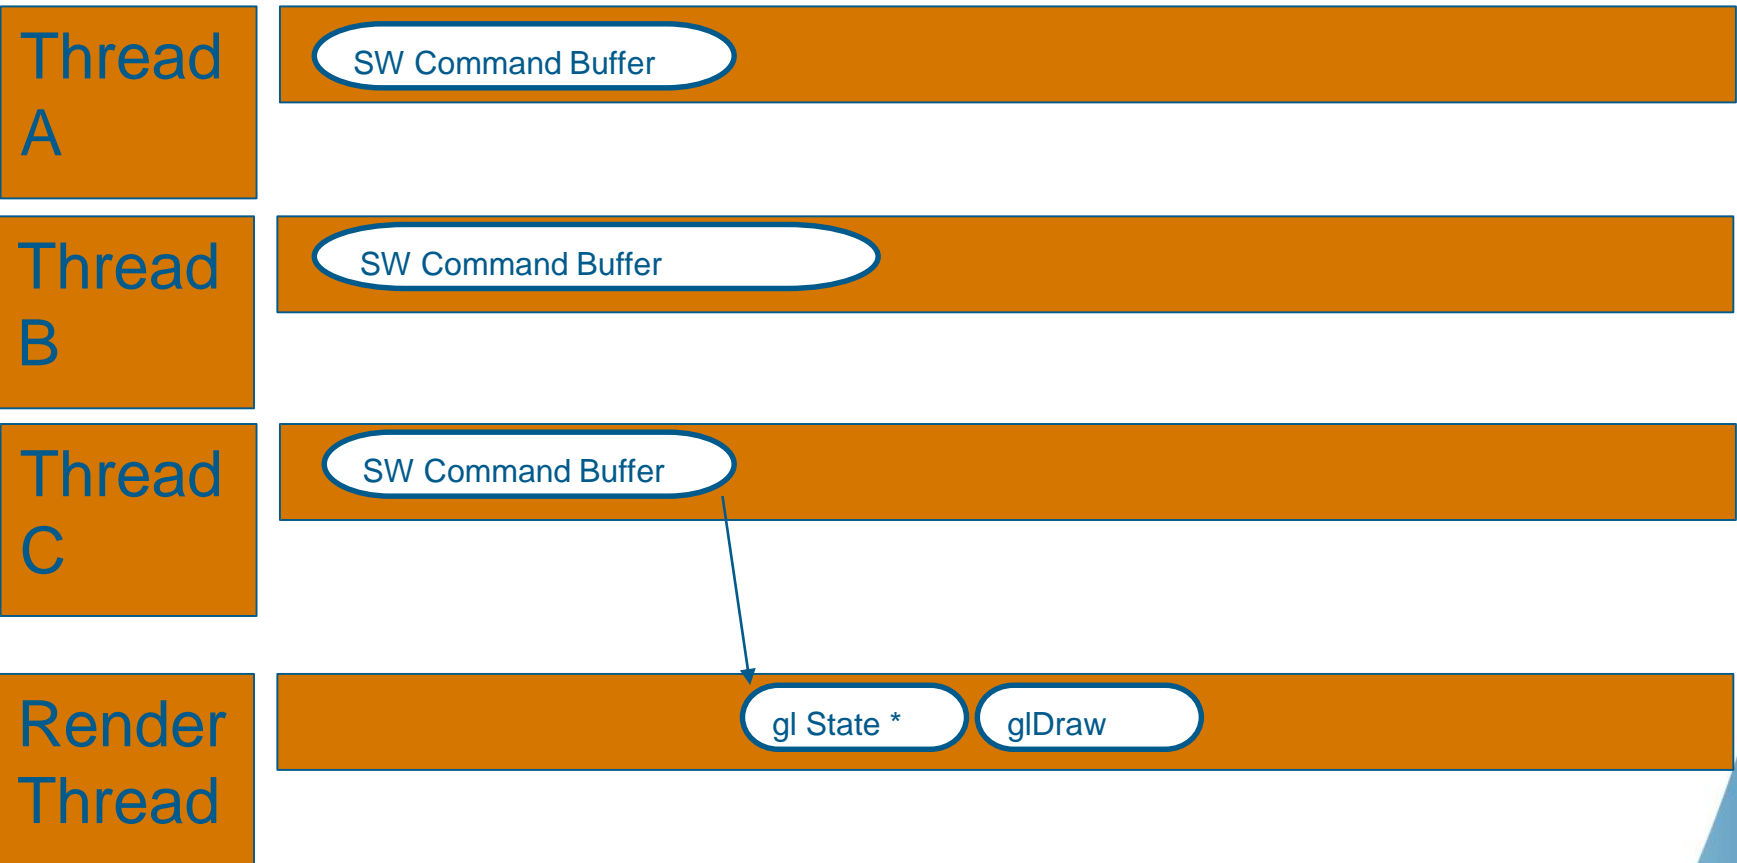
SW Command Buffer

Thread
C

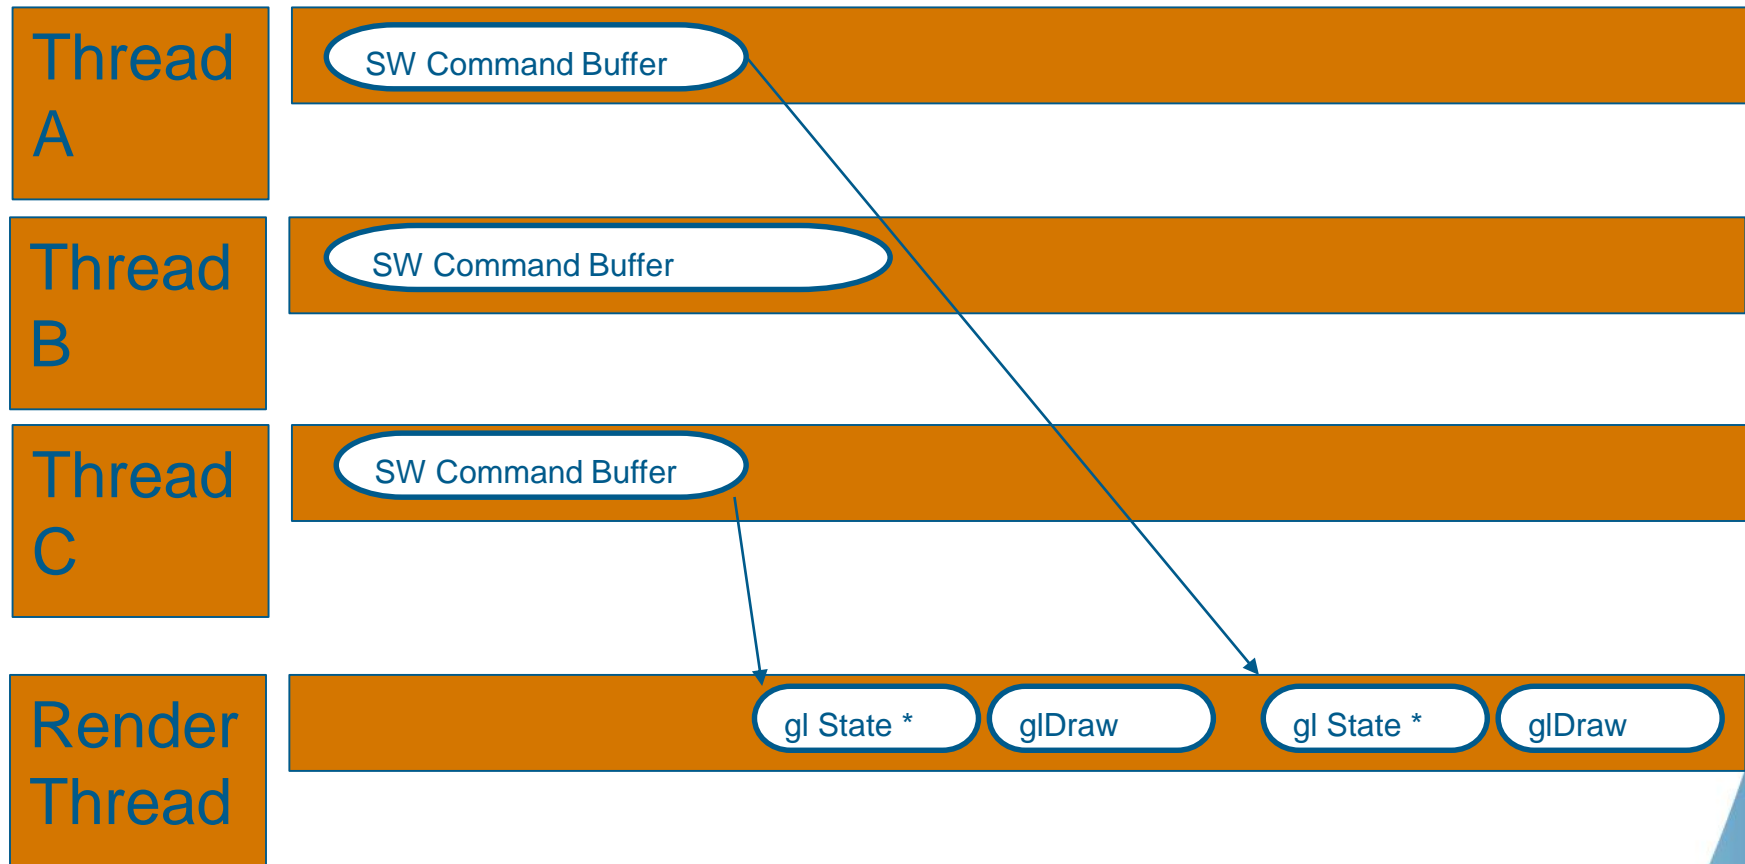
SW Command Buffer

Render
Thread

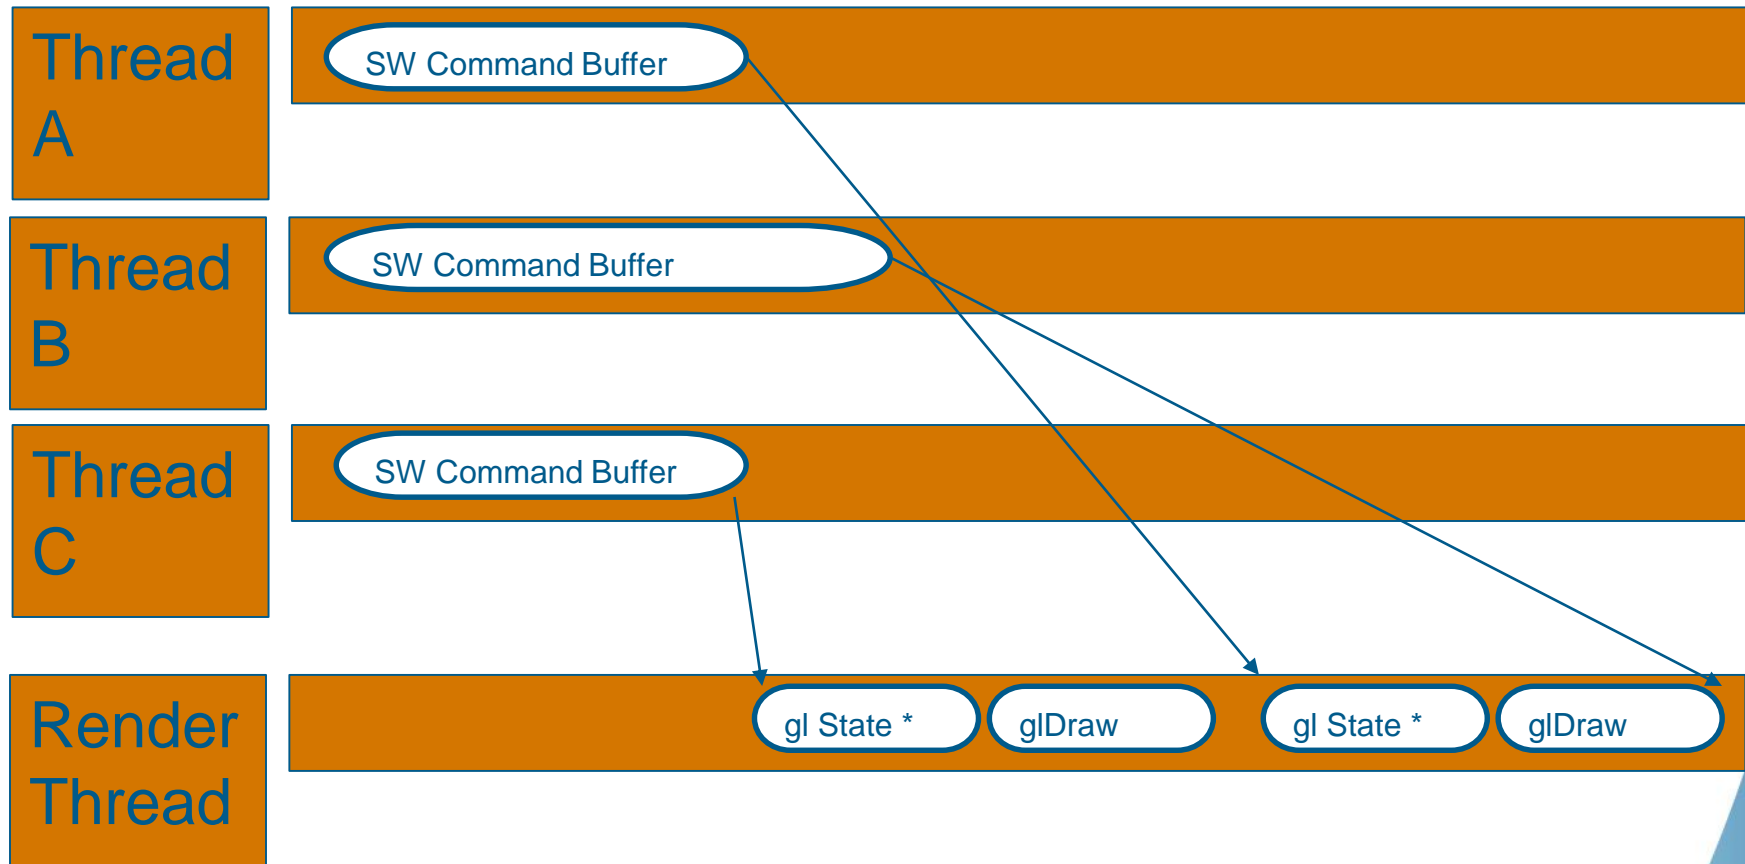
Multithreading OpenGL/OpenGLES



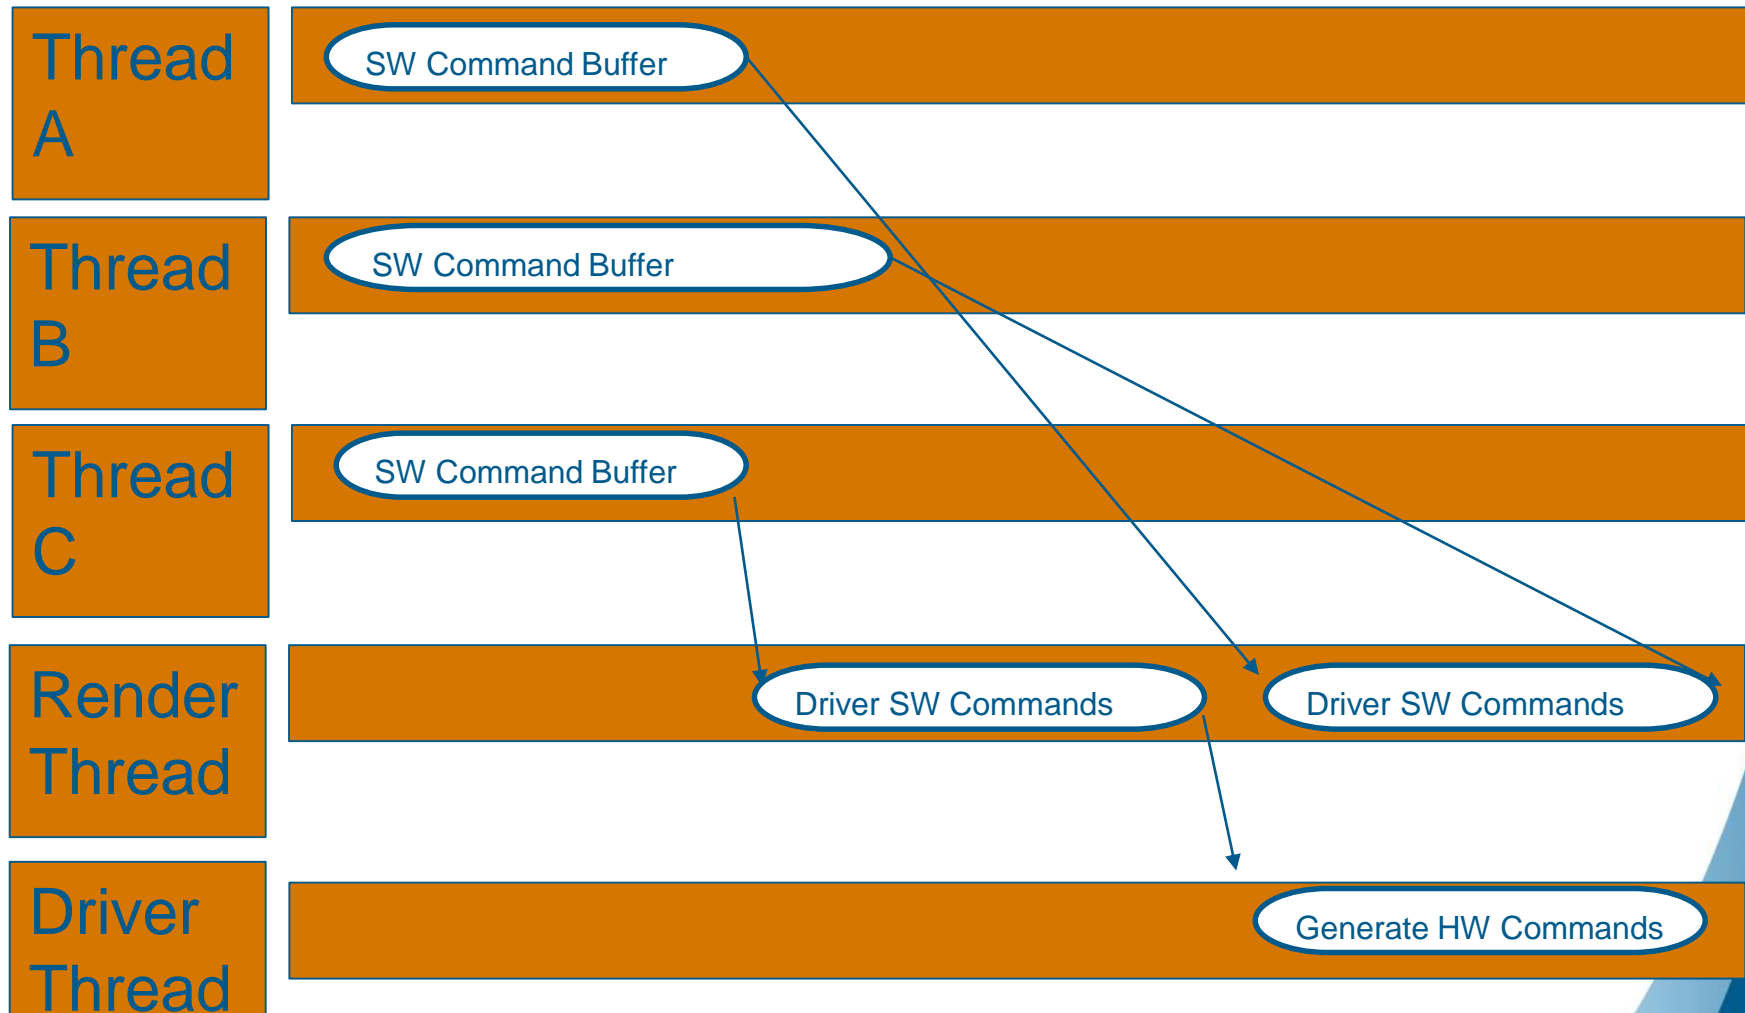
Multithreading OpenGL/OpenGLES



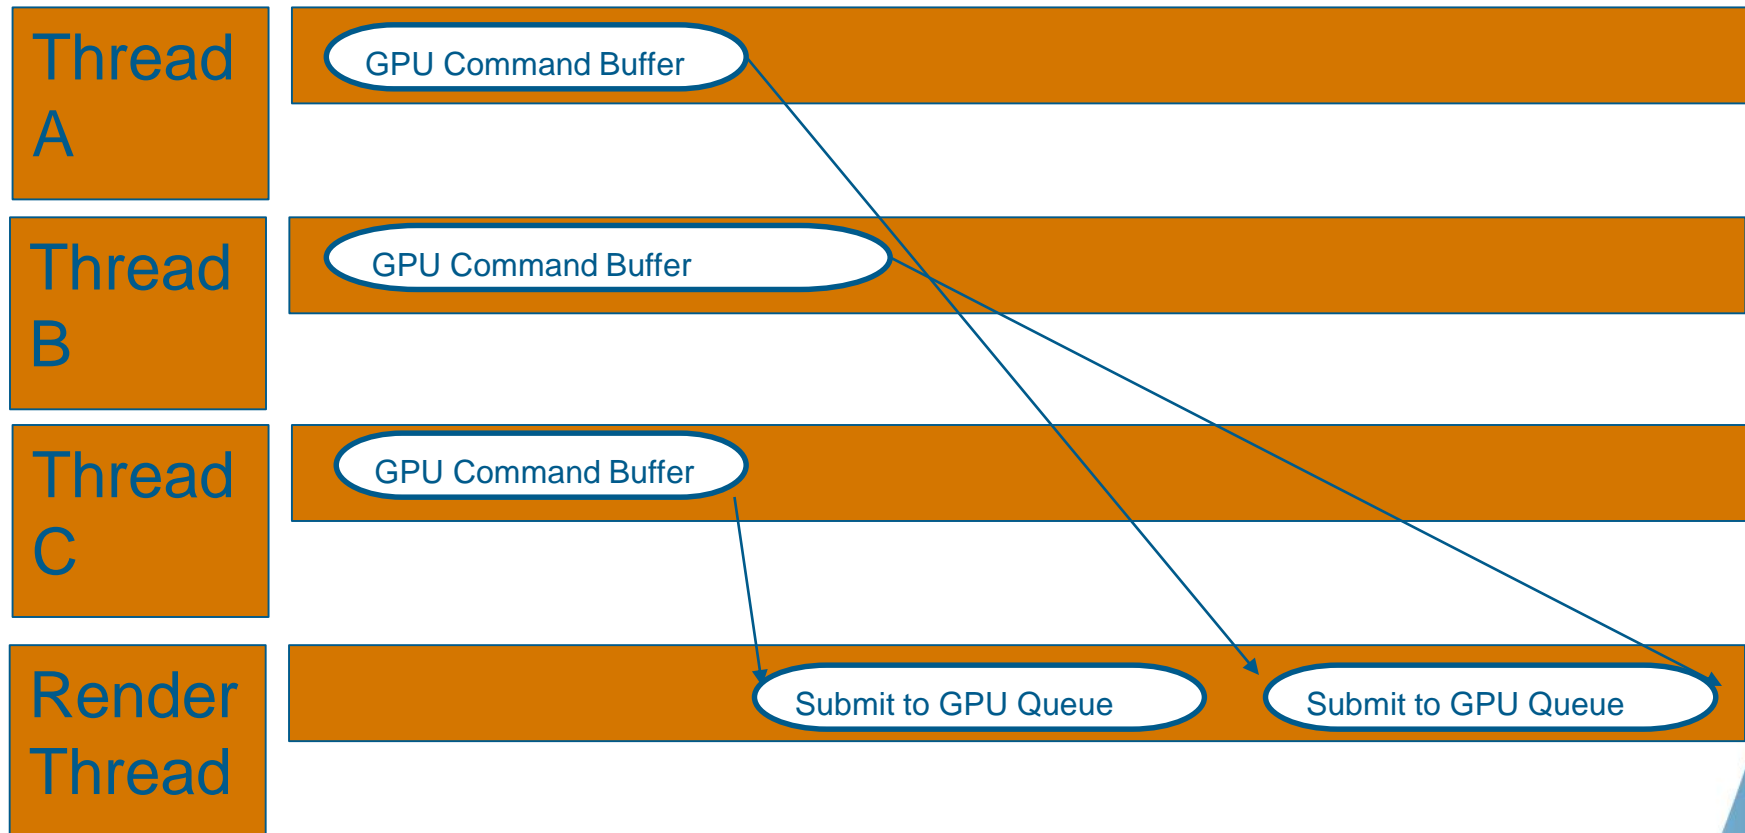
Multithreading OpenGL/OpenGLES



Multithreading OpenGL/OpenGLES



Vulkan Allows Multithreaded GPU Command Buffer Recording



OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ Lead to frame time inconsistency (“hitches”)

OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ **Lead to frame time inconsistency (“hitches”)**
- ▶ Expensive error checking
 - ▶ **Driver has to return GL_ERROR codes, expensive**

OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ **Lead to frame time inconsistency (“hitches”)**
- ▶ Expensive error checking
 - ▶ **Driver has to return GL_ERROR codes, expensive**
- ▶ Synchronization tracking
 - ▶ **When do caches flush?**
 - ▶ **When are images done being rendered to?**

OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ **Lead to frame time inconsistency (“hitches”)**
- ▶ Expensive error checking
 - ▶ **Driver has to return GL_ERROR codes, expensive**
- ▶ Synchronization tracking
 - ▶ **When do caches flush?**
 - ▶ **When are images done being rendered to?**
- ▶ Memory Allocation
 - ▶ **Where should resources get stored?**
 - ▶ **What to do if oversubscribed?**

OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ **Lead to frame time inconsistency (“hitches”)**
- ▶ Expensive error checking
 - ▶ **Driver has to return GL_ERROR codes, expensive**
- ▶ Synchronization tracking
 - ▶ **When do caches flush?**
 - ▶ **When are images done being rendered to?**
- ▶ Memory Allocation
 - ▶ **Where should resources get stored?**
 - ▶ **What to do if oversubscribed?**
- ▶ Shader compilation
 - ▶ **Each driver needs to parse GLSL/ESSL – error prone**

OpenGL/OpenGL ES Problems

- ▶ State based recompiles
 - ▶ **Lead to frame time inconsistency (“hitches”)**
- ▶ Expensive error checking
 - ▶ **Driver has to return GL_ERROR codes, expensive**
- ▶ Synchronization tracking
 - ▶ **When do caches flush?**
 - ▶ **When are images done being rendered to?**
- ▶ Memory Allocation
 - ▶ **Where should resources get stored?**
 - ▶ **What to do if oversubscribed?**
- ▶ Shader compilation
 - ▶ **Each driver needs to parse GLSL/ESSL – error prone**
- ▶ Tile Based Renderers (TBR – Mobile)
 - ▶ **Massive driver heuristics/heroics**

OpenGL/OpenGL ES Problems – Vulkan Solutions

- ▶ State based recompiles
 - ▶ **Vulkan: Pipeline State Objects**
- ▶ Expensive error checking
 - ▶ **Vulkan: Move error checking to layers**
- ▶ Synchronization tracking
 - ▶ **Vulkan: Explicit synchronization**
- ▶ Memory Allocation
 - ▶ **Vulkan: Explicit Memory Allocation**
- ▶ Shader compilation
 - ▶ **Vulkan: SPIR-V**
- ▶ Tile Based Renderers (TBR – Mobile)
 - ▶ **Vulkan: RenderPasses**

Should I use Vulkan or GL?

- ▶ Probably yes:
 - ▶ **CPU perf/high draw call throughput**
 - ▶ **Multithreaded**
 - ▶ **Frame consistency**
 - ▶ **Explicit control**
 - ▶ **Fancy GPU features (i.e. async compute)**
 - ▶ **Targeting Mobile TBRs**
- ▶ Maybe not:
 - ▶ **Fully GPU bound in shading work**
 - ▶ Vulkan may make little/no difference other than power consumption
 - ▶ **Need quick results**

Vulkan API Concepts

Vulkan Graphics API Concepts

- ▶ Command Buffer
- ▶ Queue
- ▶ Synchronization Primitives
- ▶ Shaders
- ▶ Pipeline State Object
- ▶ Descriptor Sets
- ▶ RenderPass
- ▶ Barrier
- ▶ Memory Allocation

Command Buffer

- ▶ Per-thread command recorder
- ▶ Primary or secondary
 - ▶ **Secondary for parallelizing within a renderpass (more on this later)**
- ▶ No state inheritance across boundaries
- ▶ Submit to queue

Queues

- ▶ Abstraction for hardware command queue
- ▶ vkQueueSubmit
 - ▶ **Submit command buffer(s) to GPU**
- ▶ Common case: one queue
- ▶ Other possibilities:
 - ▶ **Graphics and compute queue (async compute)**
 - ▶ **Transfer Queue**

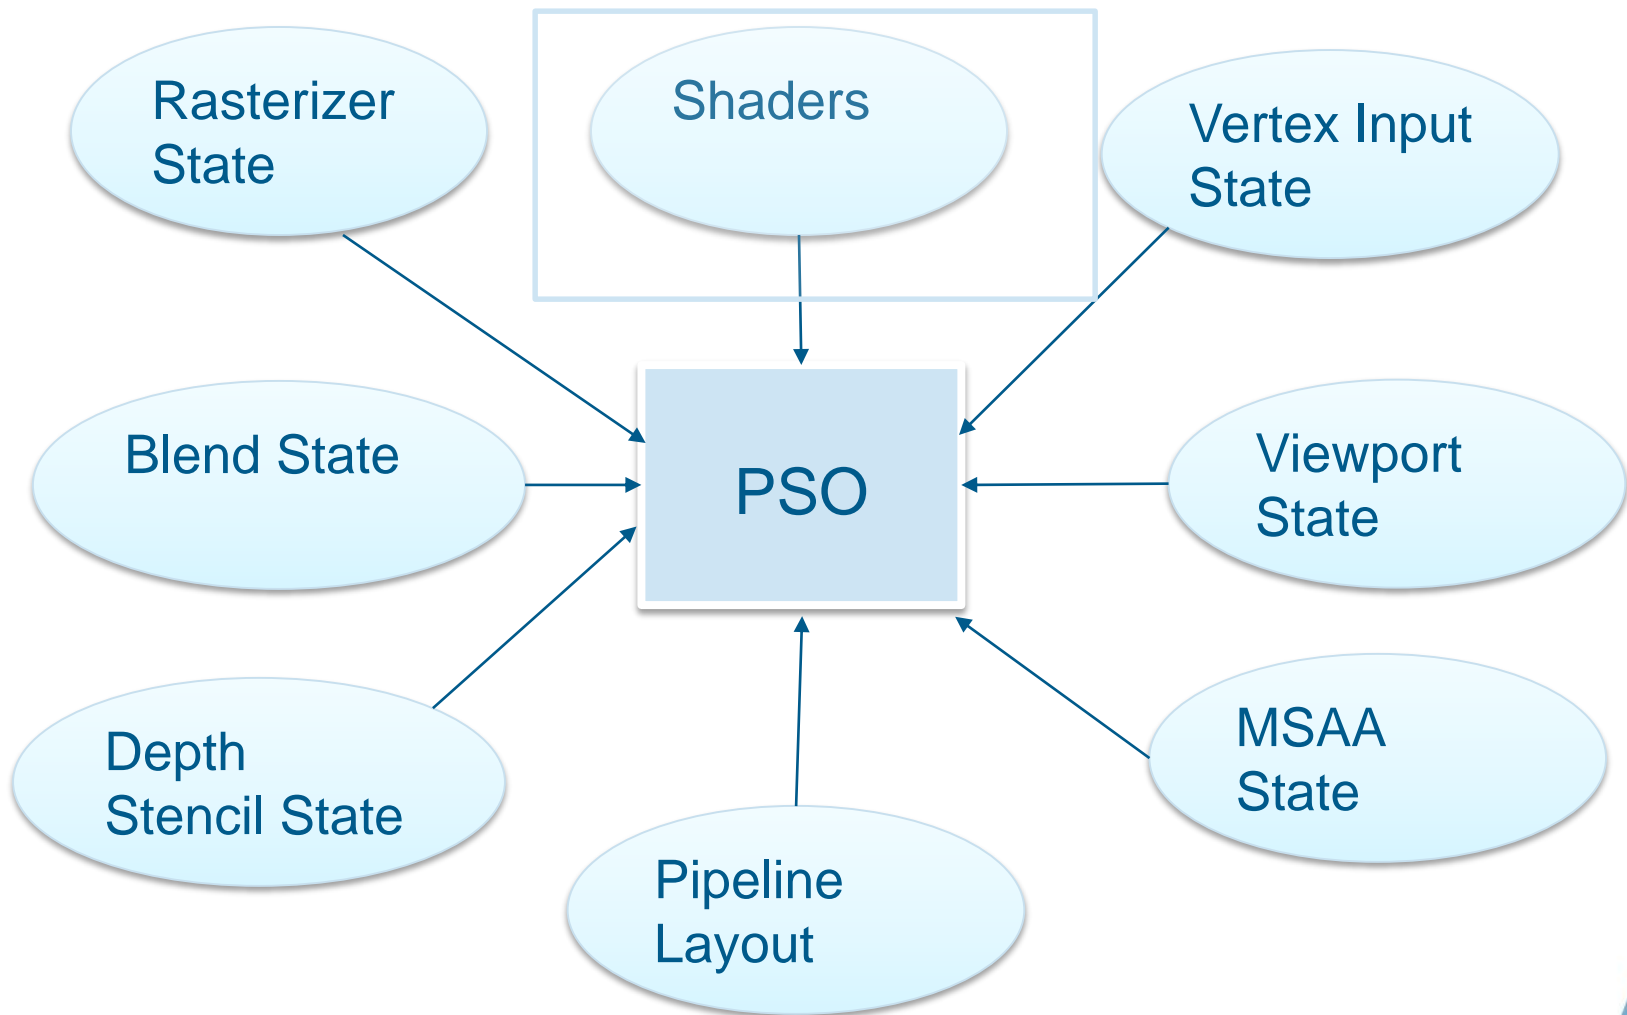
Synchronization Primitives

- ▶ Fences
 - ▶ **Determine when queue work has finished on CPU**
- ▶ Semaphore
 - ▶ **Wait for work on the GPU**
 - ▶ **Synchronize across multiple submits**
 - ▶ **Synchronize across multiple queues**
- ▶ Events
 - ▶ **Wait inside a command buffer**
 - ▶ **Signal from GPU or CPU**

Shaders

- ▶ SPIR-V
 - ▶ **Intermediate Representation (IR)**
 - ▶ **Single Static Assignment (SSA)**
- ▶ Lots of high level languages to choose from
 - ▶ **GLSL**
 - ▶ **HLSL**
 - ▶ **OpenCL C (sort of)**

Pipeline State Object (PSO)



Descriptor Sets

- ▶ Inputs to the PSO
 - ▶ **Textures**
 - ▶ **Uniforms**
 - ▶ **Storage Buffers**
 - ▶ **Etc.**
- ▶ Descriptor Set Layout
 - ▶ **Describes layout of shader inputs**
- ▶ Pipeline Layout
 - ▶ **Combines descriptor sets to define all inputs to PSO**
- ▶ Descriptor Set
 - ▶ **Populated with bindings to resources**

RenderPass

- ▶ First we should talk about tile based renderers (TBR)...

Tile Based Renderers

- ▶ Mobile GPUs have limited external memory bandwidth
 - ▶ **Slow**
 - ▶ **Uses a lot of battery**
- ▶ Mobile GPUs have fast internal memory
 - ▶ **Qualcomm Adreno calls this GMEM**
 - ▶ **ARM Mali and Apple GPUs have similar ideas**
 - ▶ **(some desktop GPUs are becoming tile based too)**
- ▶ Breaks frame up into bins or tiles

Tile Based Renderer



Tile Based Renderer



Tile Based Renderer



Rendering in two passes

- ▶ **Binning Pass**
 - ▶ **Run position portion of vertex shader to compute clip space position**
 - ▶ **Store which bin each vertex is in**
- ▶ **For each tile**
 - ▶ **Load tile with initial value**
 - ▶ LOAD – restore previous contents
 - ▶ DON'T_CARE – ignore previous contents
 - ▶ CLEAR – clear to solid color
 - ▶ **Process primitives matching that bin**
 - ▶ **Shade fragments, blend, etc.**
 - ▶ **Store tile results**
 - ▶ STORE – store to external memory
 - ▶ DON'T_CARE – throw away

Back to RenderPasses

- ▶ RenderPass defines this process
 - ▶ **What attachments are we rendering to?**
 - ▶ **What should the initial value of tile memory be?**
 - ▶ **What should we do with the results?**
- ▶ Immediate Mode Renderers
 - ▶ **This is basically BindRenderTarget/DiscardRenderTarget, but can still be useful**
- ▶ This is an overly simple explanation
 - ▶ **RenderPasses can actually represent frame graph**
 - ▶ Subpasses
 - ▶ Subpass dependencies

Barriers

- ▶ Express data dependency
- ▶ Examples
 - ▶ **Render to Texture, Bind as Shader Resource**
 - ▶ Image Barrier: COLOR_ATTACHMENT_OPTIMAL -> SHADER_READ_ONLY_OPTIMAL
 - ▶ Makes sure outstanding writes finish before reads
 - ▶ **Copy to Buffer, Read as Vertex Buffer**
 - ▶ Memory Barrier: TRANSFER_DST -> VERTEX_INPUT
- ▶ Pretty hard to get right, recommend reading this:
 - ▶ <http://themaister.net/blog/2019/08/14/yet-another-blog-explaining-vulkan-synchronization/>

Memory Allocation

- ▶ Vulkan Driver provides heaps and memory types
- ▶ Query to figure out which heaps can be used for allocation
- ▶ Allocate and bind memory to resource
- ▶ Application does all pooling
 - ▶ You have to write your own memory manager
 - ▶ Or just use VMA: <https://github.com/GPUOpen-LibrariesAndSDKs/VulkanMemoryAllocator>

Vulkan Getting Started

- ▶ Bad news:
 - ▶ **It's hard**
- ▶ Good news:
 - ▶ **There are tons of resources**
 - ▶ **It's all open source**

Vulkan Resources

- ▶ Vulkan SDK (Windows/Linux/macOS):
 - ▶ <https://vulkan.lunarg.com/>
 - ▶ **Validation Layers**
 - ▶ **Shader Compilers**
 - ▶ glslang (GLSL or HLSL SM3-5.1 -> SPIR-V)
 - ▶ DXC (HLSL SM6+)
 - ▶ **SPIRV-Tools**
 - ▶ **spirv-cross**
 - ▶ SPIR-V -> MSL/HLSL/GLSL
 - ▶ **MoltenVK (Vulkan on Metal)**
 - ▶ macOS + iOS
 - ▶ Dota Underlords shipped on it!
- ▶ What about Android?
 - ▶ **Part of NDK**
 - ▶ **Supported since Android 7**

Vulkan Repos

- ▶ glslang
 - ▶ <https://github.com/KhronosGroup/glslang>
- ▶ SPIR-V Tools
 - ▶ <https://github.com/KhronosGroup/SPIRV-Tools>
- ▶ Vulkan Validation Layers
 - ▶ <https://github.com/KhronosGroup/Vulkan-ValidationLayers>
- ▶ DirectXShaderCompiler (SM6 HLSL)
 - ▶ <https://github.com/microsoft/DirectXShaderCompiler>
- ▶ MoltenVK
 - ▶ <https://github.com/KhronosGroup/MoltenVK>

Other Useful Resources

- ▶ Vulkan Samples
 - ▶ <https://github.com/KhronosGroup/Vulkan-Samples>
- ▶ Sascha Willems Examples
 - ▶ <https://github.com/SaschaWillems/Vulkan>
- ▶ VMA (Vulkan Memory Allocator)
 - ▶ <https://github.com/GPUOpen-LibrariesAndSDKs/VulkanMemoryAllocator>
- ▶ Vulkan GPU Database
 - ▶ <http://vulkan.gpuinfo.org/>

Questions?

dginsburg@upsamplesoftware.com